

Model Armor



Google Cloud · Runtime protection

A Google Cloud inspection service for content sent to and returned by AI models. Applications submit content for a security verdict, then apply that verdict through their own integration.



AREA	CAPABILITIES AND SCOPE
Security area	Prompt and response security, sensitive-data protection and content policy.
Main features	Detects prompt injection, attempts to bypass model safeguards, harmful content and malicious URLs. Templates set filters and thresholds; sensitive-data policies can identify or redact protected information.
How it works	The API returns filter results. Supported integrations can inspect, block or sanitize the submitted content before it reaches the model or the user.
Deployment and integrations	Managed service with regional endpoints and a REST API. Supports text and document inspection; image screening is in preview and has format and integration limits.
Visibility and response	Filter findings identify triggered policies. Cloud Logging supports investigation of inspection and blocking decisions when configured.
Additional offerings	Sensitive Data Protection supplies classification and de-identification. Cloud service integrations apply Model Armor policies at supported application or gateway boundaries.
Scope and limitations	Only submitted content is inspected. It does not issue agent credentials, authorize business transactions or isolate code. Images embedded inside documents are not screened as images.

Illustrative ABB use: Screen a customer service document before the model reads it, then inspect the drafted reply for sensitive data.

Scope based on vendor documentation; the example is illustrative.

Agent 365



Microsoft · AI security and governance platforms

Microsoft's administration and governance layer for enterprise agents. It connects an agent registry and lifecycle management with identity, data protection and threat services across supported agent platforms.



AREA	CAPABILITIES AND SCOPE
Security area	Agent inventory, ownership, lifecycle governance and coordinated security administration.
Main features	Registers agents, identifies owners, maps relationships and activity, and supports lifecycle and access decisions from the Microsoft 365 administration environment.
How it works	Agent 365 coordinates oversight. Connected identity, data and threat services apply their own policies to supported agent identities, interactions and resources.
Deployment and integrations	Microsoft cloud service. Microsoft and supported third-party agents use the documented onboarding and integration paths; available capabilities depend on the connected workload.
Visibility and response	Agent maps and activity views support oversight. Security events and data-risk records feed the corresponding Microsoft security and compliance workflows.
Additional offerings	Entra Agent ID provides agent identity and access governance. Purview provides information protection and DLP. Defender adds threat detection and protection for supported agents.
Scope and limitations	Registration does not itself inspect every tool call. Coverage depends on onboarding and the security components enabled; external applications still enforce their resource permissions.

Illustrative ABB use: Maintain an accountable record of each service agent, its sponsor, access and lifecycle status.

Scope based on vendor documentation; the example is illustrative.

Amazon Bedrock AgentCore



Amazon Web Services · AI security and governance platforms

A modular AWS platform for running agents and connecting them to tools. Its security-relevant services provide identity, tool-call policy, isolated execution and activity traces.



AREA	CAPABILITIES AND SCOPE
 Security area	Agent hosting, identity, tool authorization, execution isolation and operational visibility.
 Main features	Runtime isolates sessions. Identity manages workload credentials. Gateway exposes approved APIs and tools. Policy applies fine-grained rules before gateway tool calls execute.
 How it works	The agent requests a tool through Gateway. Policy evaluates the call; configured credentials and downstream authorization determine access to the target resource.
 Deployment and integrations	Managed AWS services usable together or independently. Supports custom and open-source agent frameworks, external models, APIs and Model Context Protocol (MCP) servers.
 Visibility and response	Observability records execution traces using OpenTelemetry-compatible signals. These traces show the steps, outputs and failures of instrumented agent sessions.
 Additional offerings	Code Interpreter and Browser provide managed execution environments. Memory retains task context. Bedrock Guardrails is a separate content-screening service that complements tool policy.
 Scope and limitations	Gateway policy covers routed calls. Direct connections need separate controls. Content filtering and business approvals remain distinct from hosting, credential management and session isolation.

 **Illustrative ABB use:** Give a service agent scoped API access and a separate execution environment for diagnostic scripts.

Scope based on vendor documentation; the example is illustrative.

Prisma AIRS



Palo Alto Networks · AI security and governance platforms

Palo Alto Networks’ AI security platform combines discovery, runtime protection and model supply-chain capabilities. Its modules operate at different points in AI development and production.



AREA	CAPABILITIES AND SCOPE
 Security area	AI asset inventory, model and artifact security, application protection and agent access.
 Main features	Maps AI workloads and traffic; scans model artifacts; runs adversarial simulations; and detects threats in AI interactions through network or application integrations.
 How it works	Runtime inspection applies policy to integrated traffic paths. Supply-chain analysis examines artifacts before use. Inventory and simulation provide separate findings and context.
 Deployment and integrations	Module-specific deployment across cloud, hybrid and on-premises environments. Network inspection and code-level integrations have different coverage and operating dependencies.
 Visibility and response	Inventory, security findings and runtime events support investigation. Blocking is performed by the relevant enforcement integration, not by an inventory record alone.
 Additional offerings	AI Gateway centralizes AI-traffic policy and visibility. AI Agent Identity covers agent and MCP identity governance. AI Red Teaming and AI Supply Chain Security extend lifecycle coverage.
 Scope and limitations	The suite combines distinct capabilities. Deploying one module does not provide every other function or cover traffic that bypasses its integration.

 **Illustrative ABB use:** Combine AI inventory, model-artifact review and runtime protection across several service applications.

Scope based on vendor documentation; the example is illustrative.

AI Defense



Cisco · AI security and governance platforms

Cisco’s AI security platform combines asset discovery, supply-chain scanning, security analysis and runtime enforcement. It covers AI applications and supported agent and MCP interactions.



AREA	CAPABILITIES AND SCOPE
 Security area	AI asset visibility, model and tool supply-chain risk, application security and runtime protection.
 Main features	Discovers AI assets; scans model files, repositories and MCP servers; performs adversarial analysis; and applies security, privacy and safety policies to live traffic.
 How it works	Runtime integrations inspect prompts, responses and supported tool calls. MCP policies can allow or block tools using configured policy and asset-risk information.
 Deployment and integrations	Supports cloud and customer VPC or on-premises AI POD deployments. In VPC and AI POD configurations, sensitive traffic stays local; management metadata goes to Cisco.
 Visibility and response	Violation dashboards capture policy events. Security-console integrations, including Splunk, connect AI findings with wider incident investigation and operations.
 Additional offerings	Cloud Visibility, Supply Chain Risk Management and Model and Application Validation complement Runtime Protection. These are separate capability areas within AI Defense.
 Scope and limitations	Discovery, scanning and inline protection have different integration requirements. Coverage depends on the application and traffic path; target-system authorization remains necessary.

 **Illustrative ABB use:** Identify unapproved AI connections and apply policy to model and MCP traffic in a supported service environment.

Scope based on vendor documentation; the example is illustrative.

Falcon AI Detection and Response



CrowdStrike · Runtime protection

CrowdStrike’s AIDR service collects AI activity and applies security and data policies. For agents, its MCP proxy inspects messages exchanged with connected tools.



AREA	CAPABILITIES AND SCOPE
Security area	AI interaction visibility, prompt and data protection, and tool-traffic enforcement.
Main features	The MCP collector inspects tool discovery, requests and returned content. Guard checkpoints can allow, block or transform a message, including redacting sensitive values.
How it works	A proxy sits between the MCP client and server. It submits the relevant messages to AIDR policy checks before a tool receives them or the client receives results.
Deployment and integrations	The MCP proxy runs alongside the client and connects to AIDR. Integration is possible through client configuration, including environments without direct access to application code.
Visibility and response	Policy findings, logs and activity views support investigation. Blocking and redaction occur at the configured checkpoints on the intercepted MCP path.
Additional offerings	Other AIDR collectors support workforce browsers, endpoints, applications and AI gateways. These extend collection beyond the MCP proxy through separate integration paths.
Scope and limitations	The MCP proxy does not inspect model prompts and responses outside MCP interactions. Coverage from other collectors requires their own deployment and configuration.

Illustrative ABB use: Inspect a service agent’s tool request and returned customer data before allowing the exchange.

Scope based on vendor documentation; the example is illustrative.

AI Control Tower



ServiceNow · AI security and governance platforms

ServiceNow's AI management product connects asset discovery, governance, security and operational oversight. AI assets are linked to owners, business services and enterprise workflows.



AREA	CAPABILITIES AND SCOPE
 Security area	AI inventory, lifecycle governance, risk, operational visibility and connected security.
 Main features	Discovers and registers AI assets, records ownership and relationships, tracks usage and value, and links risk and compliance decisions to workflow records.
 How it works	Governance workflows connect asset and risk records. Supported security integrations enforce least-privilege access and block prompt injection in connected AI interactions.
 Deployment and integrations	Runs on the ServiceNow AI Platform and connects AI sources with the configuration management database (CMDB), enterprise workflows and supported external platforms.
 Visibility and response	Asset lineage, relationships and governance records support traceability. Operational and security views connect AI activity with responsible owners and follow-up actions.
 Additional offerings	Related ServiceNow offerings include Integrated Risk Management, IT Asset Management and Strategic Portfolio Management, linking AI oversight to risk, asset lifecycle and investment planning.
 Scope and limitations	Recorded approvals and risk decisions do not enforce themselves. Discovery depth and runtime action depend on the connected platform, enforcement point and workflow integration.

 **Illustrative ABB use:** Keep service-agent ownership, approved purpose, dependencies and policy exceptions in one operating record.

Scope based on vendor documentation; the example is illustrative.

AI Guardrails



Check Point · Runtime protection

Check Point’s AI Guardrails, developed from Lakera, screens content submitted by AI applications. It combines content defenses with supported controls over tool use and deviations from an agent’s assigned task.



AREA	CAPABILITIES AND SCOPE
Security area	Input and output protection, sensitive-data safeguards and agent behavior controls.
Main features	Prompt-injection detection, content moderation, data-leakage prevention and malicious-link detection. Agent defenses include a tool allow/deny list and dangerous-deviation detection.
How it works	Applications submit user input, retrieved material, outputs or supported action context for inspection. The integration uses the returned decision to block or handle a policy violation.
Deployment and integrations	API-based application integration with configurable policies. Defense settings determine which detectors screen the submitted content and action context.
Visibility and response	Policy and detector results identify flagged interactions. Runtime enforcement relies on the application or gateway applying the result at the relevant boundary.
Additional offerings	Managed guardrails combine learned detectors with rule-based filters. Policy customization and documented fine-tuning options adapt the defenses to an application’s threat profile.
Scope and limitations	Inspection cannot cover unsubmitted traffic. A content verdict is not a replacement for target-system authorization, credential lifecycle management or execution isolation.

Illustrative ABB use: Reject instructions hidden in a customer document and restrict the tools an agent can invoke.

Scope based on vendor documentation; the example is illustrative.

AI Guardrails



F5 · Runtime protection

F5's runtime security offering applies configurable guardrails to AI inputs, outputs and supported agent actions. It is integrated into the application or gateway path where a policy decision can be enforced.



AREA	CAPABILITIES AND SCOPE
Security area	AI content security, data protection and policy for tool calls and agent actions.
Main features	Screens prompts and responses, identifies disallowed or sensitive content, and applies configured restrictions to unauthorized tool calls and supported agent actions.
How it works	The integrated enforcement point evaluates an interaction against guardrail policy and applies the configured response before the interaction proceeds.
Deployment and integrations	Public-cloud, private-cloud, on-premises and air-gapped deployment options are documented. Application and gateway integrations apply guardrails to supported model and agent traffic.
Visibility and response	Logs attribute enforcement actions to guardrails. Activity views include tool calls and policy outcomes; dashboards and SIEM export support wider security operations.
Additional offerings	F5 AI Red Team is a related offering for automated adversarial analysis and vulnerability reporting. It complements live guardrails with findings about configured AI targets.
Scope and limitations	Content screening and tool-action coverage depend on the integration. Guardrails do not supply complete agent identity governance or authorize every downstream transaction.

Illustrative ABB use: Apply data and action policy between a custom service agent and its connected tools.

Scope based on vendor documentation; the example is illustrative.

AI Security Platform



Noma Security · AI security and governance platforms

Noma Security combines discovery and security posture with agent access control and runtime detection. Its platform connects configuration risk with the actions an agent takes during a task.



AREA	CAPABILITIES AND SCOPE
 Security area	Agent discovery, permissions and data relationships, access policy and behavioral protection.
 Main features	Inventories agents, models, tools and MCP servers; maps risk relationships; identifies misconfiguration; and detects injection, data exfiltration and activity outside the assigned task.
 How it works	Access policy governs approved tools and connections. Runtime detection evaluates action sequences using identity, data and session context on supported integrations.
 Deployment and integrations	Integrations cover supported employee endpoints, SaaS agent platforms and custom AI applications. Collection and enforcement depend on the connected platform and runtime.
 Visibility and response	Risk maps and session context support investigation. Documented responses include alerting, blocking, data masking and routing an action for human review.
 Additional offerings	AI-SPM supplies discovery and posture management. Access Control governs allowed connections. AI-DR supplies runtime detection and response; AI Red Teaming adds adversarial analysis.
 Scope and limitations	Capabilities differ across environments and integration paths. A discovered asset is not necessarily under inline enforcement, and unconnected systems are outside the observed scope.

 **Illustrative ABB use:** Connect a service agent’s broad customer-data access with a later attempt to send records externally.

Scope based on vendor documentation; the example is illustrative.

AI Security and Governance Platform



Zenity · AI security and governance platforms

Zenity combines agent visibility and configuration analysis with runtime policy and detection. It relates what an agent can access to what it does across a sequence of actions.



AREA	CAPABILITIES AND SCOPE
Security area	Agent discovery, security posture, exposure analysis and runtime behavior protection.
Main features	Builds an agent inventory, maps data access and permissions, identifies risky configuration and reachable attack paths, and monitors execution on supported platforms.
How it works	Discovery and posture establish context. Runtime boundaries enforce policy on connected action paths; detection and response use activity signals to identify unsafe behavior.
Deployment and integrations	Supported integrations cover business agents in SaaS platforms, custom cloud agents and endpoint or coding agents. Visibility and enforcement depth vary by platform.
Visibility and response	Activity and relationship views support investigation of tool calls, retrieved content and handoffs. Supported runtime integrations can block unsafe actions as they occur.
Additional offerings	AI Observability, AI Security Posture Management and AI Exposure Management add inventory and risk context. Runtime enforcement and AI Detection and Response address live activity.
Scope and limitations	Equivalent protection across all environments is not implied. A platform connector’s discovery capability is distinct from its support for inline blocking.

Illustrative ABB use: Identify an agent with excessive access and follow its activity across a multi-system service task.


Scope based on vendor documentation; the example is illustrative.

AI Red Teaming

Zscaler · AI security assurance



Zscaler’s AI Red Teaming automates adversarial analysis of AI models, applications and agents. It generates security findings that help explain weaknesses in a configured target.



AREA	CAPABILITIES AND SCOPE
Security area	AI security assurance, vulnerability discovery and remediation reporting.
Main features	Prebuilt and custom attacks, domain-specific datasets and multiple input formats support analysis of threats such as prompt injection, data exposure and unsafe application behavior.
How it works	The service interacts with a connected AI target, records its responses and produces findings. These findings describe the target and scenarios included in the analysis.
Deployment and integrations	Commercial cloud offering connected to supported AI application endpoints. Domain data and target configuration shape the scope of the analysis.
Visibility and response	Reports and dashboards identify findings and remediation priorities. Supported workflow integrations help track follow-up across development and security teams.
Additional offerings	Custom attack scenarios, domain-specific attack generation and continuous analysis extend the core service beyond one-off scans. Related Zscaler security services have separate runtime roles.
Scope and limitations	This product produces findings; it does not itself authorize live tool calls or act as a production blocking gateway. Results cannot establish coverage of every possible agent action.

Illustrative ABB use: Obtain structured security findings alongside the architecture, runtime-protection and governance information for an AI application.

Scope based on vendor documentation; the example is illustrative.